



# Ray support r4: training process

371 fit presentations per recorded epoch; sampled observations vary, not paired generalization evidence.


### Native build-depth Huber




### Sampled target-to-surface distance




### Sampled hard free-space intrusion




### LiDAR fallback among Actors with views



### Median parameter-gradient group norm



### Observed step time under shared resources



No checkpoint resume. Actual optimizer updates: 11130. Skipped presentations: 0.

Native loss excludes unobserved Actors. Memory contention and concurrent work affect timing; loss reduction does not establish physical accuracy.