

Validity Is Not Difficulty: Hazard-Preserving Physical 3D World Compilation for Autonomous Driving

Anonymous CVPR submission

Paper ID *****

Abstract

tee.

038

001 *Driving-world reconstruction systems often conflate two*
 002 *causes of difficulty: a scene may be hard because it con-*
 003 *tains a legitimate hazardous actor, or invalid because re-*
 004 *constructed geometry contradicts sensor evidence. Treating*
 005 *both as low confidence can obtain superficially safer worlds*
 006 *by deleting precisely the actors that matter. We introduce*
 007 *HARP-3D, a hazard-preserving physical 3D world com-*
 008 *piler that separates appearance from collision geometry*
 009 *and artifact validity from behavioral hazard. The compiler*
 010 *aligns multi-frame Actor evidence in canonical coordinates*
 011 *and exposes four surface actions—KEEP, PROJECT, COM-*
 012 *plete, and UNKNOWN—without changing Actor identity,*
 013 *trajectory, or extent. On frozen Argoverse 2, ray-certified*
 014 *compilation eliminates paired free-space violations, re-*
 015 *duces target LiDAR depth error from 0.207 to 0.154m*
 016 *and symmetric Chamfer from 0.254 to 0.175m, and re-*
 017 *tains all Actor and hazard state. A nearest-surface alterna-*
 018 *tive reaches 0.168m Chamfer but leaves 84.9% matched-*
 019 *ray violations, exposing a geometry–safety boundary. A*
 020 *nuScenes-only factorized repair-or-abstain head then cov-*
 021 *ers 75.6% of AV2 Actors zero-shot, reduces false repair from*
 022 *16.6% under always-repair to 9.0%, and obtains 0.182m*
 023 *selective Chamfer with 0.981 hazard AUROC and zero*
 024 *cross-input shift. On exact-match source Actors it rejects*
 025 *every geometrically harmful repair, yet its selected Actors*
 026 *have 2.28× larger long-tail motion error, showing why re-*
 027 *pair and trajectory authority must remain separate. Post-*
 028 *hoc attribution further assigns 49.5% of AV2 validity sensi-*
 029 *tivity to observation-window length. A pre-frozen sparsity-*
 030 *consistent recovery raises exact-once fresh nuScenes cov-*
 031 *erage from 5.7% to 26.8% but lowers repair AUROC from*
 032 *0.782 to 0.747 and is rejected. In a retained-source fixed*
 033 *lattice, separate task authority covers 49.4% of action sets*
 034 *and reduces conditional cost from 0.150 to 0.028 with-*
 035 *out changing Actor or hazard state. We therefore report*
 036 *empirical transfer while explicitly withholding a domain-*
 037 *invariant, conformal, closed-loop, or road-safety guaran-*

1. Introduction

039

040 Reconstructed driving worlds are increasingly used for per-
 041 ception analysis, counterfactual replay, and planning eval-
 042 uation. Their geometry, however, is rarely a literal collision
 043 model. A rendered Gaussian may explain image appearance
 044 while occupying free space, duplicating an Actor shell, or
 045 flickering across time. Conversely, a perfectly valid Actor
 046 may be difficult because it cuts in, brakes hard, or crosses
 047 the Ego path. A useful world compiler must distinguish
 048 these cases: *world validity is not world difficulty*.
 049

050 This distinction matters because a generic confidence fil-
 051 ter has an easy but scientifically undesirable failure mode.
 052 It can improve aggregate collision proxies by suppressing
 053 uncertain or nearby Actors. Such a world looks safer be-
 054 cause legitimate hazards disappeared, not because its phys-
 055 ical state became more faithful. We instead ask whether
 056 evidence-inconsistent geometry can be repaired while iden-
 057 tity, lifecycle, trajectory, size, time-to-collision, and hazard
 058 events remain fixed.

059 We propose HARP-3D, a physical compilation layer be-
 060 tween a reconstructed scene and downstream task queries.
 061 It represents a world as

$$\mathcal{W}_t = (\mathcal{G}_t^{\text{app}}, \mathcal{S}_t^{\text{phys}}, \mathcal{A}_t, \mathcal{P}_t, \mathbf{q}_t^{\text{val}}), \quad (1)$$

061

062 where appearance Gaussians and physical collision surfaces
 063 have separate semantics. The compiler uses multi-frame
 064 Actor alignment, ray free/occupied evidence, temporal sup-
 065 port, and source provenance to select one of four local ac-
 066 tions. Unsupported regions become UNKNOWN, never free
 067 by default, while Actors remain in the scene.

068 Our second component factorizes validity from hazard
 069 with disjoint inputs and paired counterfactual diagnostics.
 070 Changing a clean Actor from ordinary driving to a legal
 071 near-miss should not change its repair score; injecting a
 072 duplicate shell or temporal pop should not change its haz-
 073 ard score. A selective interface repairs an Actor only when

nuScenes-calibrated runtime evidence supports it and otherwise returns the clean query. Our third component links physical state quality to downstream reliability. An Actor residual distribution is projected onto a candidate trajectory boundary, yielding a continuous cost whose conditional density supports arbitrary budgets and multiple horizons.

Our contributions are:

- a hazard-preserving physical compiler with explicit KEEP/PROJECT/COMPLETE/UNKNOWN semantics and separate appearance/collision state;
- a nuScenes-only repair-or-abstain factorization protocol that measures cross-input leakage and reports its AV2 risk-coverage boundary;
- a continuous task-cost density and monotone multi-horizon reliability interface, with an exact-identity audit that keeps physical-repair and trajectory authority distinct.

2. Related Work

Driving-world reconstruction. Multi-sensor driving datasets such as nuScenes [2] and Argoverse 2 [14] provide calibrated imagery, LiDAR, ego pose, maps, and tracked Actors. LiDAR is 2.5D ray evidence whose free-space intervals are lost by an unordered point-only view [4]. Neural scene graphs model dynamic objects explicitly [9], while 3D Gaussian Splatting provides a high-quality appearance representation [5]. Our work does not reinterpret Gaussian opacity as collision occupancy. It adds a separate, queryable physical surface with sensor provenance.

Uncertainty and reliability. Deep ensembles expose useful empirical predictive disagreement [7], and conformalized quantile methods can turn residuals into calibrated marginal intervals under appropriate assumptions [10]. Interval-bound methods can certify a network over a specified input box, including for robust OOD detection [1]; they do not certify that the box models an unseen sensor domain. Driving logs violate many simple exchangeability assumptions through scene correlation and domain shift. We therefore report empirical calibration explicitly, preserve claim boundaries, and use monotone structure only for runtime consistency—not as a formal road-safety guarantee.

Cross-sensor LiDAR generalization. LiDAR domain shift includes sensor sampling as well as scene and appearance shift. Source-only subsampling consistency can improve transfer across scan densities [6], while partially accumulated geometry and sequence cues provide a complementary common representation [11]. These results motivate our frozen sparsity intervention, but do not make source consistency a certificate for actor-level repair selection; we therefore retain exact-once source and external

evaluations.

Our distinction. Prior confidence, occupancy, and reconstruction-quality signals often mix physical invalidity, visibility, and behavioral difficulty. HARP-3D instead defines separate evidence inputs, paired interventions, and failure denominators so that an improvement cannot be obtained by deleting hazardous Actors or relabeling unknown space as free.

3. Problem Formulation

Let Actor i have an immutable semantic state

$$A_i = (ID_i, X_{i,0:H}, D_i), \quad (2)$$

containing identity, trajectory, and dimensions. Appearance primitives may be associated with A_i , but collision queries consume only an Actor-local physical surface S_i^{phys} and its known/unknown mask.

We separate two random variables. Artifactness a_i depends on sensor and geometric evidence E_i , physical surface state, and provenance:

$$a_i = P(\text{artifact} \mid E_i, S_i^{\text{phys}}, P_i). \quad (3)$$

Hazardness h_i depends on Actor dynamics, map context, and a candidate Ego trajectory τ :

$$h_i = P(\text{hazard} \mid X_{i,0:H}, \tau, \mathcal{M}). \quad (4)$$

The goal is paired conditional invariance, not unconditional statistical independence. For the same clean Actor, a legal safe-to-hazard intervention should preserve a_i . For the same behavior, a clean-to-artifact intervention should preserve h_i .

For task reliability, let $n_{\tau,t}$ be the local boundary normal, $d_{i,t}(\tau)$ the signed Actor clearance, and $\epsilon_{i,t}$ the Actor residual. We define

$$C(\tau, H) = \max_{i,t \leq H} \frac{|n_{\tau,t}^\top \epsilon_{i,t}|}{\max(|d_{i,t}(\tau)|, \epsilon)}. \quad (5)$$

The runtime interface returns $R(\tau, H, b) = P(C \leq b)$ and enforces $\partial R / \partial b \geq 0$ and $\partial R / \partial H \leq 0$.

4. Method

4.1. Actor-Preserving Physical Compilation

For each Actor, we transform LiDAR endpoints, rays, and associated primitives into a canonical Actor frame,

$$x_{i,t}^{\text{actor}} = T_{\text{actor},t}^{-1} T_{\text{ego},t} T_{\text{sensor}} x_t^{\text{sensor}}. \quad (6)$$

The fused representation stores surface support, observed-free intervals, temporal counts, local normals, viewing-direction diversity, and provenance. Compilation is deterministic and local:




Figure 1. HARP-3D compiles sensor evidence into an Actor-preserving physical world, separates validity from hazard, and exposes task-conditioned reliability. Appearance Gaussians never serve directly as collision occupancy.

- KEEP: stable sensor-supported surface;
- PROJECT: a near-surface primitive contradicts observed free space and is projected to supported geometry;
- COMPLETE: a local hole has sufficient multi-frame and multi-view support;
- UNKNOWN: evidence is missing, contradictory, or insufficient.

Unknown is a first-class state and never silently becomes free. All four actions preserve $(ID_i, X_{i,0:H}, D_i)$. Appearance primitives may attach to the repaired surface using center distance, covariance-normal alignment, depth, and visibility, but physical queries do not read Gaussian opacity.

4.2. Validity-Hazard Factorization

The validity encoder receives ray contradiction, surface residual, temporal support, provenance, visibility, and local geometry. Its label is whether the frozen compiler’s target Chamfer is no worse than that of the untouched clean query. The hazard encoder receives only TTC, clearance, closing speed, braking, and crossing geometry. Our factorized model uses two structurally disjoint low-capacity heads; a shared-input two-head network is the learned baseline:

$$\mathcal{L} = \text{BCE}(s_{\text{rep}}, y_{\text{rep}}) + \text{BCE}(s_{\text{haz}}, y_{\text{haz}}). \quad (7)$$

We fit features and weights on nuScenes train only. On a disjoint calibration fold, threshold q is the maximum-coverage value satisfying the population false-repair bound

$$\hat{R}^+(q) = \frac{1 + \sum_{i=1}^n \mathbf{1}[s_i \geq q, y_i = 0]}{n + 1} \leq \alpha. \quad (8)$$

Accepted Actors use the compiled surface; abstained Actors retain the clean query. Paired input swaps measure whether either score reads the other task’s variables. The finite-sample interpretation requires exchangeability; AV2 uses the frozen q only for empirical zero-shot evaluation.

4.3. Actor Residuals and Boundary-Cost Density

For Actor i and horizon t , a low-capacity head predicts a 2D residual distribution $(\mu_{i,t}, \Sigma_{i,t})$. Analytic projection onto $n_{\tau,t}$ gives mean $n^\top \mu$ and variance $n^\top \Sigma n$. Equation 5 converts these residuals and deterministic clearance into a continuous trajectory cost.

We model $y = \log(1+C)$ with a small conditional Gaussian mixture,

$$p_\theta(y \mid z^{\text{val}}, z^{\text{haz}}, \tau, H) = \sum_k \pi_k \mathcal{N}(y; \mu_k, \sigma_k^2), \quad (9)$$

which exposes a closed-form marginal CDF for any budget. A low-rank, input-conditioned dependence model combines the frozen marginals across horizons. This factorization separates learned Actor uncertainty from analytic task geometry.

The analytic boundary also exposes deterministic geometry sensitivity. Let $a_j = |n_j^\top e_j|$ and $m_j = \max(d_j, \epsilon)$. For $C(d) = \max_j a_j / m_j$ and a uniform signed-clearance shift $d' = d + \delta$,

$$|C(d') - C(d)| \leq \max_j \frac{a_j |\delta|}{m_j m'_j} \leq \frac{a_{\max} |\delta|}{\epsilon^2}. \quad (10)$$

This follows from the max operator and clipped denominator being 1-Lipschitz; it is an error-propagation bound, not a distributional safety statement.

4.4. Monotone Runtime Surface and SceneIR

The final runtime object is a frozen grid over horizon and cost budget with monotone interpolation. Budget CDF values are non-decreasing; prefix-horizon reliability is non-increasing. Groupwise isotonic maps may calibrate predefined conditions on a disjoint source fold, but do not change these structural constraints.

The two authorities compose without implication. The repair gate first chooses the surface representation,

$$\tilde{S}_i = g_i S_i^{\text{comp}} + (1 - g_i) S_i^{\text{query}}, \quad g_i = \mathbf{1}[s_i^{\text{rep}} \geq q]. \quad (11)$$

The downstream task separately grants authority only if its multi-horizon reliability satisfies the requested budget, $a_{i,t}(c) = \mathbf{1}[R_i(t, c \mid \tilde{S}_i, \tau) \geq \rho]$. In particular, $g_i = 1 \not\Rightarrow a_{i,t}(c) = 1$: evidence can support a physical repair while Actor motion remains uncertain.

SceneIR stores Actor identities and transforms, physical surfaces, known/unknown masks, provenance, and validity fields. The AV2 adapter performs only coordinate, unit, timestamp, category, and sensor-opportunity normalization. It does not fine-tune, recalibrate, or change thresholds using AV2 quality.



Figure 2. The first metadata-frozen AV2 main case. From left: synchronized RGB with observed Actor returns, paired synthetic artifacts, the four-action ray-certified output, and motion-compensated sparse camera depth. Camera/crop selection does not read RGB appearance or method quality. This is a calibrated evidence overlay, not photorealistic reconstruction.

Table 1. nuScenes-only selective repair and AV2 zero-shot transfer. False repair is measured over all Actors; coverage is the accepted fraction. Chamfer is for the repair-or-clean-query output. Leak is the maximum mean prediction shift under a paired swap of the other task’s input. Each threshold is reused unchanged; “fresh” is the metadata-frozen exact-once nuScenes cohort.

Dataset	Head	Repair AUROC $\uparrow$	Hazard AUROC $\uparrow$	Coverage $\uparrow$	False repair $\downarrow$	Selective CD $\downarrow$	Max leak $\downarrow$
nuScenes test	shared	64.26	91.66	11.84	3.51	0.231	28.86
nuScenes test	factorized	64.91	98.11	3.95	0.44	0.246	0.00
AV2 zero-shot	shared	77.97	92.89	55.21	3.15	0.190	23.46
AV2 zero-shot	factorized	69.96	98.05	75.55	8.99	0.182	0.00
nuScenes fresh	factorized	78.23	90.60	5.69	0.00	0.214	0.00
nuScenes fresh	sparsity-consistent	74.73	90.60	26.83	2.44	0.197	0.00

5. Experiments

Data and protocol. We train and select models only on nuScenes. Argoverse 2 Sensor validation logs form a zero-shot external test. Before reading any method output, we sort the 150 validation UUIDs and select every fifth log, yielding 20 quantitative and 10 qualitative logs. No AV2 fine-tuning, post-hoc calibration, threshold selection, or failed-scene deletion is allowed. Final nuScenes and AV2 cohorts are evaluated exactly once after the method freezes.

Baselines. For physical compilation we compare naive reconstruction, a legacy validity filter, V6.7 inward-ray repair, and HARP-3D. For factorization we compare a shared encoder, a two-head shared trunk, independent encoders, and paired factorization. Reliability baselines include a scalar confidence, independent horizon classifiers, a continuous marginal density, and the full joint multi-horizon surface.

Metrics. Physical metrics include free-space violation, ray termination, point-to-surface error, LiDAR depth consistency, temporal jitter, ghost components, and surface completeness. Hazard preservation reports Actor/ID/lifecycle retention, trajectory and TTC shift, and event-count change. Factorization reports artifact/hazard AUROC and AUPRC, clean-hazard false-artifact rate,

paired prediction shift, and cross-probe leakage. Reliability reports log likelihood, integrated Brier score, calibration error, Spearman correlation, monotonic violations, and run-time.

Zero-shot Actor-local surface evidence. Following the dynamic-Actor canonicalization boundary used by NeuRAD [13], we transform ego-frame AV2 LiDAR points inside each annotated cuboid into its Actor frame. Unlike appearance-oriented seed initialization, we never mirror points: every collision surfel must retain sensor and temporal support. We fuse two of every three frames at 0.12 m and reserve the remaining frames as unseen target geometry. Across 1,046 Actors and 1.075M stable surfels, fusion reduces the mean target-to-surface distance from 0.864 m to 0.131 m and raises recall at 0.20 m from 26.16% to 85.72% (Table 2).

The paired corruption atlas contains clean, observed-free ghost, duplicate-shell, temporal-flicker, and supported-hole probes without changing Actor identity, trajectory, size, or hazard. It yields 100.0% artifact recall, 0.00% clean-hazard false-artifact rate, and 100.0% Actor/hazard retention. These perfect paired scores verify the deterministic compiler contract and input separation; they are not reported as natural-artifact generalization.




Figure 3. Frozen P7 safety envelope. Left and center: risk-coverage and geometry-coverage curves; stars mark the unchanged nuScenes-calibrated thresholds and never select an AV2 operating point. Right: Integrated Gradients relative to the nuScenes-training mean [12]. AV2 reliance on observation-frame count exposes a sensor-opportunity shortcut. Curves and attributions are descriptive rather than a cross-domain certificate.

AV2 Actor surface	Dist. (m) ↓	Recall (%) ↑
Single observed frame	0.864	26.16
Multi-frame canonical surfels	0.131	85.72

Table 2. Zero-shot held-out LiDAR support on the frozen 20-log AV2 quantitative cohort (1,046 Actors, including 241 hazardous Actors). The comparison uses real sensor points; paired synthetic corruptions are evaluated separately as a compiler contract.

Surface	Free (%) ↓	Recall (%) ↑	CD (m) ↓
Single-frame input	100.0	53.36	0.254
Nearest canonical	84.91	62.60	0.168
Ray-certified	0.00	59.56	0.175

Table 3. Frozen zero-shot AV2 quantitative result. Nearest-surface projection fits geometry best but violates matched-ray free space; ray-certified projection accepts a small Chamfer cost. Later held-out frames are target-only.

Four-action physical compilation. We next freeze one single-frame query per Actor and reserve all later held-out frames as target-only geometry. The compiler reads only the build surface and query evidence. A nearest-canonical projection initially reduces Chamfer to 0.168 m, but ray-specific evaluation reveals 84.91% early terminations in the paired observed-free interval. Our final operator is ray-certified: a primitive with direct LiDAR provenance projects to its matched observed termination; without one it becomes UNKNOWN. Across 433 quantitative Actors (147 hazardous), this reduces paired free-space violations from 100.0% to 0.00%, target depth error from 0.207 m to 0.154 m, and Chamfer from 0.254 m to 0.175 m, while raising ray-termination consistency from 56.94% to 64.38% (Table 3). Actor identity, trajectory, extent, and hazard labels are retained for 100.0% of Actors. The zero free-space rate is a paired contract result; depth, ray termination, and Chamfer use independent target-only LiDAR.

Frozen camera evidence and an exposed boundary. We project each pre-registered qualitative Actor back into synchronized AV2 ring-camera frames using the official calibration and motion compensation. Camera identity maximizes only the number of visible query LiDAR returns, with a fixed camera-order tie break; RGB is decoded afterward and every case is retained. All 30 cases have calibrated

RGB, sparse depth, four-action, and before/after video evidence; the median/minimum visible query counts are 82/14 (Fig. 2). Importantly, the frozen set exposes a limit: although mean Chamfer improves, 17.19% of all 634 Actors worsen individually. Thus P3 establishes aggregate zero-shot geometry, not a per-Actor repair certificate; this motivates a selective “repair or abstain” validity head trained and calibrated only on nuScenes.

Visibility-conditioned physical evidence. We further partition target-only LiDAR evidence into early termination, matched hit, late return, and unmatched ray; reverse surface evidence is supported, contradicted by observed free space, occluded, or off-ray unknown. Unknown is never counted as free. On all 634 Actors, query-to-compiled hit recall rises from 49.69% to 69.83%, early termination falls from 3.395% to 2.868%, and visible-surface precision rises from 99.61% to 99.77% (F-score 0.663→0.822). Yet only 64.04% of Actors add no visible violation. Provenance attribution explains the trade-off: completion supplies 95.12% of new early rays but 14.96 new hits per such ray, while only 13.49% of contradicted output points come from completion. We therefore retain the net 9,439-ray reduction, reject a universal certificate, and route Actor-level authority through abstention rather than delete completion.

nuScenes-to-AV2 selective factorization. The strict multi-frame surface contract yields 29/56/228 Actors for nuScenes train/calibration/test; we retain the small training set rather than loosen admissibility. On nuScenes test, factorization is non-inferior to the shared head for repairability (64.91 vs. 64.26 AUROC) and raises hazard AUROC from 91.66 to 98.11, while maximum paired cross-input shift falls from 28.86% to 0.00%. The conservative threshold covers 3.95% with 0.44% population false repairs.

Without AV2 fitting, the same selector covers 75.55% of 634 Actors and reduces false repair from 16.56% under always-repair to 8.99%. Clean-query/always/selective Chamfer is 0.258/0.177/0.182 m. Thus abstention costs only 5 mm mean Chamfer relative to always-repair while nearly halving false repairs. Hazard AUROC is 98.05, hazard coverage is 92.17%, and factorized cross-input shift remains zero. However, coverage changes from 8.93% on nuScenes calibration to 75.55% on AV2; Table 1 therefore supports empirical transfer, not a cross-domain conformal guarantee.

Interpretable safety envelope. We retain the model, standardizer, thresholds, and all Actor rows and evaluate 21 fixed coverage levels without refitting or recalibration. The factorized score shifts from mean/median .761/.945 on nuScenes calibration to .978/.999992 on AV2 (Wasserstein .217, KS .704), explaining the operating-point coverage jump. Integrated Gradients [12] relative to the nuScenes-training mean attributes 18.25% of absolute validity sensitivity to observation-frame count on nuScenes test but 49.53% on AV2 (Fig. 3). Structural cross-input derivatives remain exactly zero: task factorization prevents hazard leakage, but does not make validity inputs sensor invariant. We therefore preserve the empirical P4 result while rejecting the stronger interpretation of a domain-invariant selector.

We then propagate fixed 0.10 nuScenes-training-standard-deviation feature boxes through the frozen validity MLP in FP64. Only 33.33% of nominal nuScenes-test selections are robust to the full box, versus 91.23% on consumed AV2. Yet 10.76% of those AV2 robust selections are target false repairs: local threshold invariance can stably preserve a domain-shifted error. Sensor-opportunity perturbations create mean logit widths of 1.550/1.285 on nuScenes/AV2, and query-point count is the largest single-feature interval source for 83.33/95.11% of Actors. This actor-level certificate makes the shortcut inspectable, while explicitly separating network stability from repair correctness.

Fresh exact-once selector boundary. Before this read, we exclude every P4 train/calibration/test scene and freeze 20 of the remaining 106 locally available scenes by equally

Table 4. Frozen physical-trajectory interface on exact-match nuScenes test Actors. Cost and error are row-level 90th percentiles over retained V6.7 outcomes; FS/flip gives event counts. Physical selection is not trajectory authority.

P4 group	Actors	q90 cost ↓	q90 error ↓	FS / flip ↓
selected	5	0.283	1.577	0 / 0
abstained	83	0.102	0.693	7 / 47
geometric harm	23	0.076	0.543	0 / 4

spaced official scene index. The 20 scenes yield 123 Actors and are read once, with no replacement or model/threshold update. Motivated by source-only LiDAR sparsity consistency [6], a pre-frozen candidate increases coverage from 5.69% to 26.83% and improves selective Chamfer from 0.214 to 0.197 m. It nevertheless lowers repair AUROC from 78.23 to 74.73 and fails the pre-registered 2-point non-inferiority margin (Table 1). A descriptive common-coverage audit follows the selective-classification AURC formulation of Zhou *et al.* [15]: candidate/P4 AURC is 0.126/0.106 (lower is better), confirming global ranking loss rather than only threshold mismatch. We retain the P4 selector; the recovery remains a negative ablation.

Physical repair is not trajectory authority. We exact-join P4 identities to retained V6.7 multi-horizon outcome rows through the official nuScenes scene order and DriveStudio instance tokens. Only 5 P4-train Actors from two scenes align, so we reject joint fitting rather than recycle calibration/test Actors. The frozen descriptive audit instead contains 88 test Actors and 39,285 rows (Table 4). P4 selects 5 Actors and abstains on all 23 whose compiled Chamfer is worse than the query; selected-and-harmful is therefore zero. Its 1,781 selected rows contain no false-safe or decision-flip events. Nevertheless, selected q90 motion error is 1.577 m versus 0.693 m for abstained Actors (2.28×), reaching 4.246 m at 3 s. The physical gate protects representation quality but cannot replace the downstream reliability authority. These are retained source outcomes, not a causal effect of repair or an execution of the reliability model.

We also stress the analytic geometry-to-cost bound on 575,596 retained profiles under 6 fixed $\pm 0.05/0.10/0.20$ m shifts. FP64 evaluation has zero violations. At -0.20 m, full-source mean/q99 absolute cost shift is 0.0292/0.1450, while only 0.841% of rows cross signed clearance. P4-selected versus abstained mean shift is 0.00085/0.00456 (5.38× more stable) with no selected crossings. Selected Actors can therefore be motion-difficult yet geometrically stable, further separating the three axes of repairability, geometry sensitivity, and trajectory uncertainty.

Table 5. Retained-source 2×2 composition on exact-match P5 Actors. Authority coverage is the accepted fraction of fixed two-action sets; cost and unsafe rate are measured over accepted sets (all sets for “none”). Identical B0/B1 and B2/B3 action metrics are the intended physical/task non-interference contract, not a causal planning effect.

		Surface CD ↓	Authority cov. (%) ↑	Actual cost ↓	Unsafe (%) ↓	Actor/hazard retention
B0	query + none	0.229	100.00	0.1497	16.45	88/32
B1	HARP-3D + none	0.223	100.00	0.1497	16.45	88/32
B2	query + P346	0.229	49.43	0.0279	1.15	88/32
B3	HARP-3D + P346	0.223	49.43	0.0279	1.15	88/32

Minimal composed-authority utility. Task-relevant monitoring should propagate prediction error to the queried planning cost [3], rather than treat all prediction errors alike. We therefore form a retained-source 2×2 audit over query/HARP-3D surfaces and no-task/frozen-P346 authority. The 88 exact Actors (32 hazardous) yield 1,228 fixed two-action sets at the P346 artifact’s held-out 2.5 s horizon. Frozen 90% requested reliability authorizes 49.43% of sets, reducing conditional mean actual cost from 0.1497 to 0.0279 (81.37%) and unsafe-set rate from 16.45% to 1.15%. Independently, selective HARP-3D lowers mean surface Chamfer from 0.229 to 0.223 m with all Actors and hazards retained (Table 5). Physical branches have identical action outcomes by construction: this validates interface composition, not closed-loop benefit from repair.

Qualitative protocol. Main-paper and supplement scenes are frozen from metadata and registered failure strata before rendering method outputs. Panels show RGB, LiDAR, physical surface, ray evidence, compiler action, and the repaired world. We report all frozen cases, including failures, rather than selecting only visually favorable examples.

6. Limitations

HARP-3D is not a formal guarantee of real-road safety. Sparse or one-sided sensing can leave large regions unknown, and deterministic completion is intentionally conservative. Actor cuboids and tracked identities may themselves be wrong. Visibility-conditioned evidence improves in aggregate, but only 64.0% of Actors add no observed-ray violation; completion exchanges a small early-return tail for about 15 times as many new matched hits. We therefore neither claim a universal surface certificate nor delete completion post hoc. The selective model has only 29 admissible nuScenes training Actors; its calibration coverage changes from 8.9% on nuScenes to 75.6% on AV2, directly precluding a cross-domain exchangeability claim. Its population false-repair bound is not conditional risk among selected Actors. Attribution identifies a concrete sensor-opportunity shortcut: observation-frame count accounts for 49.5% of AV2 validity sensitivity. This is model sensitivity rather than causality. Likewise, a network-decision in-

terval uses standardized stress boxes rather than calibrated sensor noise; 10.76% of robustly selected AV2 Actors remain false repairs, so stable predictions are not correctness certificates. A source-only sparsity-consistent recovery increases the frozen operating-point coverage on 20 fresh nuScenes scenes but fails global AUROC and AURC ranking, so we reject rather than promote it. Exact identity alignment supplies only two P4-train scenes and five Actors, so we do not jointly fit the physical and reliability heads. Zero false-safe events among five selected test Actors is descriptive, not a confidence bound. Likewise, the geometry-to-cost inequality uses uniform clearance shifts, not a measured sensor-error distribution or closed-loop intervention. The composed-authority audit reuses source outcomes, conditions on abstention, and does not execute vehicle dynamics; open-loop planning metrics can be biased by non-perceptual cues [8]. It therefore demonstrates interface utility, not collision or policy improvement. Interface, representation, dynamics, and calibration shift otherwise remain unresolved. The current physical compiler targets Actor-local surfaces; static topology and deformable Actors remain incomplete. Finally, preservation of identity, trajectory, and size guarantees only hazard functions that depend on those variables; appearance-dependent human judgments may still change.

7. Conclusion

We formulate driving-world compilation around a simple boundary: validity is not difficulty. HARP-3D repairs evidence-inconsistent physical geometry while preserving legitimate hazardous Actors and factorizes repairability from hazard with explicit abstention. Frozen nuScenes-to-AV2 transfer reduces false repair without using hazard inputs or deleting hazardous Actors. Its audits turn abstract guarantee gaps into measurable boundaries: cross-domain validity relies strongly on sensor opportunity, locally robust network decisions can remain wrong, a sparsity-consistent recovery improves coverage but fails fresh exact-once ranking, physically admissible repairs concentrate on Actors with wider long-horizon motion error, and those Actors can nevertheless be less sensitive to geometric clearance perturbations. A retained-source fixed-lattice audit further shows

that physical surface choice and task authority can compose without changing Actor, hazard, or action denominators. SceneIR therefore keeps repairability, geometry sensitivity, and multi-horizon task authority distinct, alongside appearance, collision state, unknown space, and provenance. We retain the simpler factorized selector and report rejected or non-causal boundaries rather than expanding the demonstrated claim.

References

- [1] Julian Bitterwolf, Alexander Meinke, and Matthias Hein. Certifiably adversarially robust detection of out-of-distribution data. In *Advances in Neural Information Processing Systems*, 2020. 2
- [2] Holger Caesar, Varun Bankiti, Alex H. Lang, Sourabh Vora, Venice Erin Liong, Qiang Xu, Anush Krishnan, Yu Pan, Giancarlo Baldan, and Oscar Beijbom. nuscenes: A multi-modal dataset for autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020. 2
- [3] Alec Farid, Sushant Veer, Boris Ivanovic, Karen Leung, and Marco Pavone. Task-relevant failure detection for trajectory predictors in autonomous vehicles. In *Proceedings of the 6th Conference on Robot Learning*, pages 1959–1969, 2023. 7
- [4] Peiyun Hu, Jason Zigar, David Held, and Deva Ramanan. What you see is what you get: Exploiting visibility for 3d object detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11001–11009, 2020. 2
- [5] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkuehler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics*, 42(4), 2023. 2
- [6] Hyeonseong Kim, Yoonsu Kang, Changgyoon Oh, and Kuk-Jin Yoon. Single domain generalization for lidar semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17587–17598, 2023. 2, 6
- [7] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, 2017. 2
- [8] Zhiqi Li, Zhiding Yu, Shiyi Lan, Jiahao Li, Jan Kautz, Tong Lu, and Jose M. Alvarez. Is ego status all you need for open-loop end-to-end autonomous driving? In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14864–14873, 2024. 7
- [9] Julian Ost, Fahim Mannan, Nils Thuerey, Julian Knodt, and Felix Heide. Neural scene graphs for dynamic scenes. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021. 2
- [10] Yaniv Romano, Evan Patterson, and Emmanuel J. Candès. Conformalized quantile regression. In *Advances in Neural Information Processing Systems*, 2019. 2
- [11] Jules Sanchez, Jean-Emmanuel Deschaud, and François Goulette. Domain generalization of 3d semantic segmentation in autonomous driving. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 18077–18087, 2023. 2
- [12] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 3319–3328, 2017. 5, 6
- [13] Adam Tonderski, Carl Lindström, Georg Hess, William Ljungbergh, Lennart Svensson, and Christoffer Petersson. Neurad: Neural rendering for autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. 4
- [14] Benjamin Wilson, William Qi, Tanmay Agarwal, John Lambert, Jagjeet Singh, Siddhesh Khandelwal, Bowen Pan, Ratnesh Kumar, Andrew Hartnett, Jhony Kaesemodel Pontes, Deva Ramanan, Peter Carr, and James Hays. Argoverse 2: Next generation datasets for self-driving perception and forecasting. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021. 2
- [15] Han Zhou, Jordy Van Landeghem, Teodora Popordanoska, and Matthew B. Blaschko. A novel characterization of the population area under the risk coverage curve (aurc) and rates of finite sample estimators. In *Proceedings of the 42nd International Conference on Machine Learning*, pages 79226–79253, 2025. 6