

Validity Is Not Difficulty: Hazard-Preserving Physical 3D World Compilation for Autonomous Driving

Anonymous CVPR submission

Paper ID *****

Abstract

road-safety guarantee.

038

001 *Driving-world reconstruction systems often conflate two*
 002 *causes of difficulty: a scene may be hard because it con-*
 003 *tains a legitimate hazardous actor, or invalid because re-*
 004 *constructed geometry contradicts sensor evidence. Treating*
 005 *both as low confidence can obtain superficially safer worlds*
 006 *by deleting precisely the actors that matter. We introduce*
 007 *HARP-3D, a hazard-preserving physical 3D world com-*
 008 *piler that separates appearance from collision geometry*
 009 *and artifact validity from behavioral hazard. The compiler*
 010 *aligns multi-frame Actor evidence in canonical coordinates*
 011 *and exposes four surface actions—KEEP, PROJECT, COM-*
 012 *plete, and UNKNOWN—without changing Actor identity,*
 013 *trajectory, or extent. On frozen Argoverse 2, ray-certified*
 014 *compilation eliminates paired free-space violations, re-*
 015 *duces target LiDAR depth error from 0.207 to 0.154 m and*
 016 *symmetric Chamfer from 0.254 to 0.175 m, and retains all*
 017 *Actor and hazard state. A target-nearest proximity diagnos-*
 018 *tic raises hit recall from 49.7% to 69.8%, but only 64.0%*
 019 *of Actors add no visible violation, exposing an aggregate-*
 020 *to-Actor boundary. A nuScenes-only factorized repair-or-*
 021 *abstain head then covers 75.6% of AV2 Actors zero-shot,*
 022 *reduces false repair from 16.6% under always-repair to*
 023 *9.0%, and obtains 0.182 m selective Chamfer with 0.981*
 024 *hazard AUROC and zero cross-input shift. On exact-match*
 025 *source Actors it rejects every geometrically harmful repair,*
 026 *yet its selected Actors have 2.28× larger long-tail motion*
 027 *error, showing why repair and trajectory authority must re-*
 028 *main separate. Literal first-return audits expose 8.742%*
 029 *source exposure and independently confirm 14.817% on*
 030 *fresh AV2—8.29× its target-nearest proxy—while proving*
 031 *a sharp boundary: deleting points cannot create earlier*
 032 *returns, yet every frozen deletion policy worsens Cham-*
 033 *fer. In a retained-source fixed lattice, separate task au-*
 034 *thority covers 49.4% of action sets and reduces conditional*
 035 *cost from 0.150 to 0.028 without changing Actor or hazard*
 036 *state. We therefore report empirical transfer while explicitly*
 037 *withholding a domain-invariant, conformal, closed-loop, or*

1. Introduction

039

040 Reconstructed driving worlds are increasingly used for per-
 041 ception analysis, counterfactual replay, and planning eval-
 042 uation. Their geometry, however, is rarely a literal collision
 043 model. A rendered Gaussian may explain image appearance
 044 while occupying free space, duplicating an Actor shell, or
 045 flickering across time. Conversely, a perfectly valid Actor
 046 may be difficult because it cuts in, brakes hard, or crosses
 047 the Ego path. A useful world compiler must distinguish
 048 these cases: *world validity is not world difficulty*.
 049

050 This distinction matters because a generic confidence fil-
 051 ter has an easy but scientifically undesirable failure mode.
 052 It can improve aggregate collision proxies by suppressing
 053 uncertain or nearby Actors. Such a world looks safer be-
 054 cause legitimate hazards disappeared, not because its phys-
 055 ical state became more faithful. We instead ask whether
 056 evidence-inconsistent geometry can be repaired while iden-
 057 tity, lifecycle, trajectory, size, time-to-collision, and hazard
 058 events remain fixed.

059 We propose HARP-3D, a physical compilation layer be-
 060 tween a reconstructed scene and downstream task queries.
 061 It represents a world as

$$\mathcal{W}_t = (\mathcal{G}_t^{\text{app}}, \mathcal{S}_t^{\text{phys}}, \mathcal{A}_t, \mathcal{P}_t, \mathbf{q}_t^{\text{val}}), \quad (1)$$

061

062 where appearance Gaussians and physical collision surfaces
 063 have separate semantics. The compiler uses multi-frame
 064 Actor alignment, ray free/occupied evidence, temporal sup-
 065 port, and source provenance to select one of four local ac-
 066 tions. Unsupported regions become UNKNOWN, never free
 067 by default, while Actors remain in the scene.

068 Our second component factorizes validity from hazard
 069 with disjoint inputs and paired counterfactual diagnostics.
 070 Changing a clean Actor from ordinary driving to a legal
 071 near-miss should not change its repair score; injecting a
 072 duplicate shell or temporal pop should not change its haz-
 073 ard score. A selective interface repairs an Actor only when

nuScenes-calibrated runtime evidence supports it and otherwise returns the clean query. Our third component links physical state quality to downstream reliability. An Actor residual distribution is projected onto a candidate trajectory boundary, yielding a continuous cost whose conditional density supports arbitrary budgets and multiple horizons.

Our contributions are:

- a hazard-preserving physical compiler with explicit KEEP/PROJECT/COMPLETE/UNKNOWN semantics, separate appearance/collision state, and a literal first-return audit with an exact deletion-monotonicity boundary;
- a nuScenes-only repair-or-abstain factorization protocol that measures cross-input leakage and reports its AV2 risk–coverage boundary;
- a continuous task-cost density and monotone multi-horizon reliability interface, with an exact-identity audit that keeps physical-repair and trajectory authority distinct.

2. Related Work

Driving-world reconstruction. Multi-sensor driving datasets such as nuScenes [2] and Argoverse 2 [19] provide calibrated imagery, LiDAR, ego pose, maps, and tracked Actors. LiDAR is 2.5D ray evidence whose free-space intervals are lost by an unordered point-only view [5]. Neural scene graphs model dynamic objects explicitly [12], while 3D Gaussian Splatting provides a high-quality appearance representation [8]. Our work does not reinterpret Gaussian opacity as collision occupancy. It adds a separate, queryable physical surface with sensor provenance. Ray-potential reconstruction treats visibility as a viewing-ray constraint [15]. Evidential occupancy work renders the first predicted occupied intersection back to raw LiDAR depth [6, 7]. We distinguish that literal first return from a target-nearest proximity proxy and state exactly which operator owns each audit.

Uncertainty and reliability. Deep ensembles expose useful empirical predictive disagreement [10], and conformalized quantile methods can turn residuals into calibrated marginal intervals under appropriate assumptions [13]. Interval-bound methods can certify a network over a specified input box, including for robust OOD detection [1]; they do not certify that the box models an unseen sensor domain. Driving logs violate many simple exchangeability assumptions through scene correlation and domain shift. We therefore report empirical calibration explicitly, preserve claim boundaries, and use monotone structure only for runtime consistency—not as a formal road-safety guarantee.

Selective prediction couples conditional risk to coverage [4], while learning to defer evaluates the downstream fallback as part of the system [11]. Moreover, abstention can improve aggregate regression yet worsen a prede-

defined subgroup [16]. We therefore report both selected and fallback-composite outcomes and decompose them over the frozen physical hazard stratum; this is descriptive accounting, not subgroup calibration.

Cross-sensor LiDAR generalization. LiDAR domain shift includes sensor sampling as well as scene and appearance shift. Source-only subsampling consistency can improve transfer across scan densities [9], while partially accumulated geometry and sequence cues provide a complementary common representation [14]. These results motivate our frozen sparsity intervention, but do not make source consistency a certificate for actor-level repair selection; we therefore retain exact-once source and external evaluations.

Our distinction. Prior confidence, occupancy, and reconstruction-quality signals often mix physical invalidity, visibility, and behavioral difficulty. HARP-3D instead defines separate evidence inputs, paired interventions, and failure denominators so that an improvement cannot be obtained by deleting hazardous Actors or relabeling unknown space as free.

3. Problem Formulation

Let Actor i have an immutable semantic state

$$A_i = (ID_i, X_{i,0:H}, D_i), \quad (2)$$

containing identity, trajectory, and dimensions. Appearance primitives may be associated with A_i , but collision queries consume only an Actor-local physical surface S_i^{phys} and its known/unknown mask.

We separate two random variables. Artifactness a_i depends on sensor and geometric evidence E_i , physical surface state, and provenance:

$$a_i = P(\text{artifact} \mid E_i, S_i^{\text{phys}}, P_i). \quad (3)$$

Hazardness h_i depends on Actor dynamics, map context, and a candidate Ego trajectory τ :

$$h_i = P(\text{hazard} \mid X_{i,0:H}, \tau, \mathcal{M}). \quad (4)$$

The goal is paired conditional invariance, not unconditional statistical independence. For the same clean Actor, a legal safe-to-hazard intervention should preserve a_i . For the same behavior, a clean-to-artifact intervention should preserve h_i .

For task reliability, let $n_{\tau,t}$ be the local boundary normal, $d_{i,t}(\tau)$ the signed Actor clearance, and $\epsilon_{i,t}$ the Actor residual. We define

$$C(\tau, H) = \max_{i,t \leq H} \frac{|n_{\tau,t}^\top \epsilon_{i,t}|}{\max(|d_{i,t}(\tau)|, \epsilon)}. \quad (5)$$




Figure 1. Validity is not difficulty. Top: a metadata-locked 2×2 teaser built from the first non-hazardous (q01-a0) and hazardous (q00-a0) frozen main cases. Within each Actor pair, only the paired synthetic artifact changes; identity, trajectory, extent, camera, and hazard label remain fixed. Bottom: HARP-3D compiles sensor evidence into an Actor-preserving physical world before granting separate task authority. These are calibrated evidence overlays, not photorealistic reconstructions.

The runtime interface returns $R(\tau, H, b) = P(C \leq b)$ and enforces $\partial R / \partial b \geq 0$ and $\partial R / \partial H \leq 0$.

4. Method

4.1. Actor-Preserving Physical Compilation

For each Actor, we transform LiDAR endpoints, rays, and associated primitives into a canonical Actor frame,

$$x_{i,t}^{\text{actor}} = T_{\text{actor},t}^{-1} T_{\text{ego},t} T_{\text{sensor}} x_t^{\text{sensor}}. \quad (6)$$

The fused representation stores surface support, observed-free intervals, temporal counts, local normals, viewing-direction diversity, and provenance. Compilation is deterministic and local:

- KEEP: stable sensor-supported surface;
- PROJECT: a near-surface primitive contradicts observed free space and is projected to supported geometry;
- COMPLETE: a local hole has sufficient multi-frame and multi-view support;
- UNKNOWN: evidence is missing, contradictory, or insufficient.

Unknown is a first-class state and never silently becomes free. All four actions preserve $(ID_i, X_{i,0:H}, D_i)$. Appearance primitives may attach to the repaired surface using center distance, covariance-normal alignment, depth, and visibility, but physical queries do not read Gaussian opacity.

The physical certificate is visibility-conditioned. For a held-out ray $r = (o, v, d)$ and surface S , the target-facing state $\zeta(r, S)$ is one of early, hit, late, or unmatched. In the reverse direction, an output point's state $\eta(x, r)$ is contradicted, supported, occluded, or off-ray unknown. Only the contradicted state places x in observed free space; occluded and off-ray regions remain non-falsified rather than becoming occupied or free. This bidirectional partition makes

the failure object inspectable while explicitly withholding a completeness certificate for unseen surface.

For the literal safety audit, define the first valid positive depth within the frozen beam tube as

$$d_S(r) = \min\{s > 0 : o + sv \in S_e\}, \quad \min \emptyset = \infty. \quad (7)$$

If a policy only deletes points, $S' \subseteq S$ implies $d_{S'}(r) \geq d_S(r)$. Hence the early predicate $\mathbf{1}[d_S(r) < d^* - \tau]$ is monotone non-increasing under deletion, as is new-early relative to a fixed query surface. Hit retention and symmetric Chamfer are not monotone under this order: deletion may expose or remove a hit, and it can move the two directed Chamfer terms in opposite directions. This is a geometric certificate for one predicate, not a collision guarantee.

4.2. Validity-Hazard Factorization

The validity encoder receives ray contradiction, surface residual, temporal support, provenance, visibility, and local geometry. Its label is whether the frozen compiler's target Chamfer is no worse than that of the untouched clean query. The hazard encoder receives only TTC, clearance, closing speed, braking, and crossing geometry. Our factorized model uses two structurally disjoint low-capacity heads; a shared-input two-head network is the learned baseline:

$$\mathcal{L} = \text{BCE}(s_{\text{rep}}, y_{\text{rep}}) + \text{BCE}(s_{\text{haz}}, y_{\text{haz}}). \quad (8)$$

We fit features and weights on nuScenes train only. On a disjoint calibration fold, threshold q is the maximum-coverage value satisfying the population false-repair bound

$$\hat{R}^+(q) = \frac{1 + \sum_{i=1}^n \mathbf{1}[s_i \geq q, y_i = 0]}{n + 1} \leq \alpha. \quad (9)$$

Accepted Actors use the compiled surface; abstained Actors retain the clean query. Paired input swaps measure whether either score reads the other task’s variables. The finite-sample interpretation requires exchangeability; AV2 uses the frozen q only for empirical zero-shot evaluation.

For a frozen stratum z , selector s_i , query/compiled Chamfer (q_i, c_i) , and introduced visible-failure indicator v_i , the complete defer-to-query system admits an exact accounting decomposition:

$$\begin{aligned} G_z &= \mathbb{E}_z[s_i(q_i - c_i)] = C_z \mathbb{E}_z[q_i - c_i \mid s_i = 1], \\ M_z &= \mathbb{E}_z[s_i v_i] = C_z \mathbb{E}_z[v_i \mid s_i = 1]. \end{aligned} \quad (10)$$

where $C_z = \mathbb{E}_z[s_i]$ is repair coverage. Population gain and introduced-failure mass are the stratum-share-weighted sums of G_z and M_z . Hence Actor-state preservation, conditional risk, and complete-system utility are distinct claims.

4.3. Actor Residuals and Boundary-Cost Density

For Actor i and horizon t , a low-capacity head predicts a 2D residual distribution $(\mu_{i,t}, \Sigma_{i,t})$. Analytic projection onto $n_{\tau,t}$ gives mean $n^\top \mu$ and variance $n^\top \Sigma n$. Equation 5 converts these residuals and deterministic clearance into a continuous trajectory cost.

We model $y = \log(1+C)$ with a small conditional Gaussian mixture,

$$p_\theta(y \mid z^{\text{val}}, z^{\text{haz}}, \tau, H) = \sum_k \pi_k \mathcal{N}(y; \mu_k, \sigma_k^2), \quad (11)$$

which exposes a closed-form marginal CDF for any budget. A low-rank, input-conditioned dependence model combines the frozen marginals across horizons. This factorization separates learned Actor uncertainty from analytic task geometry.

The analytic boundary also exposes deterministic geometry sensitivity. Let $a_j = |n_j^\top e_j|$ and $m_j = \max(d_j, \epsilon)$. For $C(d) = \max_j a_j/m_j$ and a uniform signed-clearance shift $d' = d + \delta$,

$$|C(d') - C(d)| \leq \max_j \frac{a_j |\delta|}{m_j m'_j} \leq \frac{a_{\max} |\delta|}{\epsilon^2}. \quad (12)$$

This follows from the max operator and clipped denominator being 1-Lipschitz; it is an error-propagation bound, not a distributional safety statement.

4.4. Monotone Runtime Surface and SceneIR

The final runtime object is a frozen grid over horizon and cost budget with monotone interpolation. Budget CDF values are non-decreasing; prefix-horizon reliability is non-increasing. Groupwise isotonic maps may calibrate predefined conditions on a disjoint source fold, but do not change these structural constraints.

The two authorities compose without implication. The repair gate first chooses the surface representation,

$$\tilde{S}_i = g_i S_i^{\text{comp}} + (1 - g_i) S_i^{\text{query}}, \quad g_i = \mathbf{1}[s_i^{\text{rep}} \geq q]. \quad (13)$$

The downstream task separately grants authority only if its multi-horizon reliability satisfies the requested budget, $a_{i,t}(c) = \mathbf{1}[R_i(t, c \mid \tilde{S}_i, \tau) \geq \rho]$. In particular, $g_i = 1 \not\Rightarrow a_{i,t}(c) = 1$: evidence can support a physical repair while Actor motion remains uncertain.

SceneIR stores Actor identities and transforms, physical surfaces, known/unknown masks, provenance, and validity fields. The AV2 adapter performs only coordinate, unit, timestamp, category, and sensor-opportunity normalization. It does not fine-tune, recalibrate, or change thresholds using AV2 quality.

5. Experiments

Data and protocol. We train and select models only on nuScenes. Argoverse 2 Sensor validation logs form a zero-shot external test. Before reading any method output, we sort the 150 validation UUIDs and select every fifth log, yielding 20 quantitative and 10 qualitative logs. No AV2 fine-tuning, post-hoc calibration, threshold selection, or failed-scene deletion is allowed. Final nuScenes and AV2 cohorts are evaluated exactly once after the method freezes.

Baselines. For physical compilation we compare naive reconstruction, a legacy validity filter, V6.7 inward-ray repair, and HARP-3D. For factorization we compare a shared encoder, a two-head shared trunk, independent encoders, and paired factorization. Reliability baselines include a scalar confidence, independent horizon classifiers, a continuous marginal density, and the full joint multi-horizon surface.

Metrics. Physical metrics include free-space violation, ray termination, point-to-surface error, LiDAR depth consistency, temporal jitter, ghost components, and surface completeness. Hazard preservation reports Actor/ID/lifecycle retention, trajectory and TTC shift, and event-count change. Factorization reports artifact/hazard AUROC and AUPRC, clean-hazard false-artifact rate, paired prediction shift, and cross-probe leakage. Reliability reports log likelihood, integrated Brier score, calibration error, Spearman correlation, monotonic violations, and runtime.

Zero-shot Actor-local surface evidence. Following the dynamic-Actor canonicalization boundary used by NeuRAD [18], we transform ego-frame AV2 LiDAR points inside each annotated cuboid into its Actor frame. Unlike appearance-oriented seed initialization, we never mir-

Table 1. nuScenes-only selective repair and AV2 zero-shot transfer. False repair is measured over all Actors; coverage is the accepted fraction. Chamfer is for the repair-or-clean-query output. Leak is the maximum mean prediction shift under a paired swap of the other task’s input. Each threshold is reused unchanged; “fresh” denotes independently metadata-frozen exact-once cohorts.

Dataset	Head	Repair AUROC $\uparrow$	Hazard AUROC $\uparrow$	Coverage $\uparrow$	False repair $\downarrow$	Selective CD $\downarrow$	Max leak $\downarrow$
nuScenes test	shared	64.26	91.66	11.84	3.51	0.231	28.86
nuScenes test	factorized	64.91	98.11	3.95	0.44	0.246	0.00
AV2 zero-shot	shared	77.97	92.89	55.21	3.15	0.190	23.46
AV2 zero-shot	factorized	69.96	98.05	75.55	8.99	0.182	0.00
nuScenes fresh	factorized	78.23	90.60	5.69	0.00	0.214	0.00
nuScenes fresh	sparsity-consistent	74.73	90.60	26.83	2.44	0.197	0.00
AV2 fresh	factorized	65.48	98.50	77.25	11.28	0.184	0.00
AV2 fresh	sparsity-consistent	67.62	98.50	83.94	12.43	0.178	0.00




Figure 2. Physical hard evidence and boundaries. Left: target-nearest proximity undercounts literal first-return exposure on source, consumed AV2 (AV2-C), and disjoint exact-once AV2 (AV2-F). Center: deletion certifies non-increasing early returns, but every frozen policy loses hits and worsens Chamfer. Right: attribution localizes the AV2 sensor-opportunity shortcut. These are measurement, structural-direction, and model-sensitivity results—not collision or road-safety certificates.

AV2 Actor surface	Dist. (m) $\downarrow$	Recall (%) $\uparrow$
Single observed frame	0.864	26.16
Multi-frame canonical surfels	0.131	85.72

Table 2. Zero-shot held-out LiDAR support on the frozen 20-log AV2 quantitative cohort (1,046 Actors, including 241 hazardous Actors). The comparison uses real sensor points; paired synthetic corruptions are evaluated separately as a compiler contract.

ror points: every collision surfel must retain sensor and temporal support. We fuse two of every three frames at 0.12 m and reserve the remaining frames as unseen target geometry. Across 1,046 Actors and 1.075M stable surfels, fusion reduces the mean target-to-surface distance from 0.864 m to 0.131 m and raises recall at 0.20 m from 26.16% to 85.72% (Table 2).

The paired corruption atlas contains clean, observed-free ghost, duplicate-shell, temporal-flicker, and supported-hole probes without changing Actor identity, trajectory, size, or hazard. It yields 100.0% artifact recall, 0.00% clean-hazard false-artifact rate, and 100.0% Actor/hazard retention. These perfect paired scores verify the deterministic compiler contract and input separation; they are not reported as natural-artifact generalization.

Four-action physical compilation. We next freeze one single-frame query per Actor and reserve all later held-out frames as target-only geometry. The compiler reads only the build surface and query evidence. A nearest-canonical projection initially reduces Chamfer to 0.168 m, but ray-specific evaluation reveals 84.91% early terminations in the paired observed-free interval. Our final operator is ray-certified: a primitive with direct LiDAR provenance projects to its matched observed termination; without one it becomes UNKNOWN. Across 433 quantitative Actors (147 hazardous), this reduces paired free-space violations from 100.0% to 0.00%, target depth error from 0.207 m to 0.154 m, and Chamfer from 0.254 m to 0.175 m, while raising ray-termination consistency from 56.94% to 64.38% (Table 3). Actor identity, trajectory, extent, and hazard labels are retained for 100.0% of Actors. The zero free-space rate is a paired contract result; depth, ray termination, and Chamfer use independent target-only LiDAR.

Frozen camera evidence and an exposed boundary. We project each pre-registered qualitative Actor back into synchronized AV2 ring-camera frames using the official calibration and motion compensation. Camera identity maximizes only the number of visible query LiDAR returns, with a fixed camera-order tie break; RGB is decoded after-

Surface	Free (%) ↓	Recall (%) ↑	CD (m) ↓
Single-frame input	100.0	53.36	0.254
Nearest canonical	84.91	62.60	0.168
Ray-certified	0.00	59.56	0.175

Table 3. Frozen zero-shot AV2 quantitative result. Nearest-surface projection fits geometry best but violates matched-ray free space; ray-certified projection accepts a small Chamfer cost. Later held-out frames are target-only.

ward and every case is retained. All 30 cases have calibrated RGB, sparse depth, four-action, and before/after video evidence; the median/minimum visible query counts are 82/14 (Fig. 1). Importantly, the frozen set exposes a limit: although mean Chamfer improves, 17.19% of all 634 Actors worsen individually. Thus P3 establishes aggregate zero-shot geometry, not a per-Actor repair certificate; this motivates a selective “repair or abstain” validity head trained and calibrated only on nuScenes.

Visibility-conditioned physical evidence. We further partition target-only LiDAR evidence with a target-nearest proximity operator into early, matched hit, late, and unmatched; reverse surface evidence is supported, contradicted by observed free space, occluded, or off-ray unknown. Unknown is never counted as free. On all 634 Actors, query-to-compiled hit recall rises from 49.69% to 69.83%, early termination falls from 3.395% to 2.868%, and visible-surface precision rises from 99.61% to 99.77% (F-score 0.663→0.822). Yet only 64.04% of Actors add no visible violation. An independent 20-log read reproduces hit 49.72→68.99%, early 3.448→2.931%, and F-score 0.663→0.815, while only 62.72% add no visible violation and 19.50% worsen in Chamfer. Provenance attribution explains the trade-off: completion supplies 95.12% of new early rays but 14.96 new hits per such ray, while only 13.49% of contradicted output points come from completion. We therefore retain the net 9,439-ray reduction, reject a universal certificate, and route Actor-level authority through abstention rather than delete completion. These frozen values localize errors around each target return; they are not literal first-return rates, which are audited below over every candidate surface point in the beam tube.

nuScenes-to-AV2 selective factorization. The strict multi-frame surface contract yields 29/56/228 Actors for nuScenes train/calibration/test; we retain the small training set rather than loosen admissibility. On nuScenes test, factorization is non-inferior to the shared head for repairability (64.91 vs. 64.26 AUROC) and raises hazard AUROC from 91.66 to 98.11, while maximum paired cross-input shift falls from 28.86% to 0.00%. The conservative thresh-

old covers 3.95% with 0.44% population false repairs.

Without AV2 fitting, the same selector covers 75.55% of 634 Actors and reduces false repair from 16.56% under always-repair to 8.99%. Clean-query/always/selective Chamfer is 0.258/0.177/0.182 m. Thus abstention costs only 5 mm mean Chamfer relative to always-repair while nearly halving false repairs. Hazard AUROC is 98.05, hazard coverage is 92.17%, and factorized cross-input shift remains zero. However, coverage changes from 8.93% on nuScenes calibration to 75.55% on AV2; Table 1 therefore supports empirical transfer, not a cross-domain conformal guarantee.

Interpretable safety envelope. We retain the model, standardizer, thresholds, and all Actor rows and evaluate 21 fixed coverage levels without refitting or recalibration. The factorized score shifts from mean/median .761/.945 on nuScenes calibration to .978/.999992 on AV2 (Wasserstein .217, KS .704), explaining the operating-point coverage jump. Integrated Gradients [17] relative to the nuScenes-training mean attributes 18.25% of absolute validity sensitivity to observation-frame count on nuScenes test but 49.53% on AV2 (Fig. 2, right). The complete fixed-threshold risk-coverage and geometry-coverage curves remain in the supplement. Structural cross-input derivatives remain exactly zero: task factorization prevents hazard leakage, but does not make validity inputs sensor invariant. We therefore preserve the empirical P4 result while rejecting the stronger interpretation of a domain-invariant selector.

We then propagate fixed 0.10 nuScenes-training-standard-deviation feature boxes through the frozen validity MLP in FP64. Only 33.33% of nominal nuScenes-test selections are robust to the full box, versus 91.23% on consumed AV2. Yet 10.76% of those AV2 robust selections are target false repairs: local threshold invariance can stably preserve a domain-shifted error. Sensor-opportunity perturbations create mean logit widths of 1.550/1.285 on nuScenes/AV2, and query-point count is the largest single-feature interval source for 83.33/95.11% of Actors. This actor-level certificate makes the shortcut inspectable, while explicitly separating network stability from repair correctness.

Fresh exact-once selector boundary. Before this read, we exclude every P4 train/calibration/test scene and freeze 20 of the remaining 106 locally available scenes by equally spaced official scene index. The 20 scenes yield 123 Actors and are read once, with no replacement or model/threshold update. Motivated by source-only LiDAR sparsity consistency [9], a pre-frozen candidate increases coverage from 5.69% to 26.83% and improves selective Chamfer from 0.214 to 0.197 m. It nevertheless lowers repair AUROC

from 78.23 to 74.73 and fails the pre-registered 2-point non-inferiority margin (Table 1). A descriptive common-coverage audit follows the selective-classification AURC formulation of Zhou *et al.* [20]: candidate/P4 AURC is 0.126/0.106 (lower is better), confirming global ranking loss rather than only threshold mismatch. We retain the P4 selector; the recovery remains a negative ablation.

On 523 exact-once AV2 Actors, the order reverses: candidate/P4 AUROC is 67.62/65.48 and Chamfer 0.178/0.184 m, but false repair is 12.43/11.28%. External gates pass, yet this source/external reversal is cohort-dependent rather than universal; P4 remains primary.

Physical authority is not a visibility certificate. We exact-join the same 523 fresh AV2 Actors to the independently frozen bidirectional visibility audit, without changing either score or threshold. P4 selects 77.25%; its Chamfer-worsening rate falls from 19.50% under always-repair to 14.60%. Selected visible-failure risk changes only from 37.28% to 36.39%, and its one-sided 95% Wilson upper bound 40.40% remains above the always-repair point risk. Safe-visible AUROC is 0.533, and abstention captures 24.62% of visible failures. Thus P4 controls its Chamfer target but does not confidence-separate visibility risk. Requiring no COMPLETE output and conjoining with P4 lowers visible failure to 20.16% at 23.71% coverage, but Chamfer worsening rises to 43.55% and hazard coverage collapses to 3.52%; observed provenance is therefore visibility evidence, not jointly valid repair authority.

A dedicated nuScenes-only visibility head confirms this mismatch rather than solving it. Safe-visible AUROC is 0.753 on source test and 0.625 on AV2. Visibility-only selection reaches 11.63% AV2 visible risk at 8.22% coverage; conjunction with P4 reaches 10.26% (95% upper 20.98%) at 7.46% coverage. Yet dual Chamfer-worsening is 35.90% and hazard coverage is 3.52%. We reject the head: low target-specific risk obtained by losing geometry and hazardous Actors is not a safety authority.

The complete defer-to-query system preserves 100.0% of Actors and separates conditional risk from population utility. P4 has 28.11% population introduced-visible failure with 0.083 m composite Chamfer gain; P6-C has 30.78% and 0.089 m, remaining an AV2-only frontier point because its fresh-nuScenes ranking reverses. P4-and-visibility has only 0.76% introduced failure but 0.0012 m gain. In contrast, the provenance policy has 4.78% introduced failure and -0.0009 m gain, so query-only strictly dominates it. Low selected risk alone therefore does not establish fallback-system utility.

Exact stratification exposes where this utility is obtained. Hazardous Actors are 27.15% of the cohort; P4 selected visible risk is 47.37%/31.00% on hazardous/clear Actors versus 46.48%/33.86% under always-repair. Only 3 of its 48

removed failures are hazardous, yet they carry 42.86% of introduced failures and 56.42% of gain. Fresh target-nearest action attribution localizes the mismatch: hazardous/clear proxy new-early rates are 1.912%/0.980% ($1.951\times$); COMPLETE causes 94.70% of new early returns and 99.94% of new hits, whereas KEEP causes 95.68% of surface contradictions. P4's hazardous rate is 1.922% ($1.006\times$ always), and P6-C is $1.001\times$. State preservation therefore does not imply hazard-stratum risk reduction or suppression of its completion-driven mechanism.

Literal first-return correction and safety-surface boundary. The preceding frozen action audit uses the target-nearest point as a proximity attribution. Following raw-ray occupancy evaluation [6], we separately re-audit fixed source and consumed AV2 surfaces with Eq. 7. The source total/hazard proxy rates 0.966/1.436% rise to 5.956/8.742%. P17/P17R/P19 reduce hazard exposure to 7.777/8.230/8.662%, but each worsens Chamfer. Thus the certified early direction has nonzero surface cost.

Consumed AV2 first shows hazard exposure $1.912\rightarrow 13.926\%$ ($7.29\times$). A disjoint exact-once 10-log read then independently confirms all/hazard $1.370/1.788\rightarrow 10.801/14.817\%$ ($7.89/8.29\times$). All 70,220 fresh literal new-early rays are COMPLETE; completion supplies 98.398% of 86,883 new hits, only 1.237 hits/early. Figure 2 (left and center) visualizes both findings. They transfer metric failure and provenance, not model, policy, collision, or road-safety validity.

Frozen continuous and joint reliability. We inherit the V6.7 reliability objects without V7 refitting. On 10 previously unread nuScenes logs, the five-component log-cost density reduces mean integrated Brier and marginal calibration error by 28.48% and 69.38% relative to the frozen monotone-CDF control. A conditional four-horizon copula improves joint-event Brier by 16.97% on source-held-out trajectories and by 17.52% on an independently frozen 10-log cohort; calibration-error reductions are 71.85% and 53.37% (Table 4). The downstream P346 authority reaches 32.84% q90 coverage at 9.09% unsafe rate on that consumed fresh cohort, but source held-out-horizon risk rises to 26.71%. We therefore treat density, dependence, and task authority as separate inherited evidence, not as a V7 repair effect or a formal cross-domain calibration claim.

Physical repair is not trajectory authority. We exact-join P4 identities to retained V6.7 multi-horizon outcome rows through the official nuScenes scene order and DriveStudio instance tokens. Only 5 P4-train Actors from two scenes align, so we reject joint fitting rather than recycle calibration/test Actors. The frozen descriptive audit instead

Table 4. Frozen V6.7 reliability evidence inherited without V7 refitting. P183 has 10 fresh logs; source held-out has 3,742 trajectories; P201 has 1,846 trajectories from 10 fresh logs. $\Delta B/\Delta E$ are Brier/calibration-error reductions. P346 reuses P201 and remains development evidence.

Object	Split	$\Delta B \uparrow$	$\Delta E \uparrow$	cov./unsafe $\uparrow\downarrow$
density	P183 fresh	28.48	69.38	—
joint	source held-out	16.97	71.85	—
joint	P201 fresh	17.52	53.37	—
P346	P201 reused	—	—	32.84/9.09
P346	source held-out <i>H</i>	—	—	29.42/26.71

Table 5. Frozen physical-trajectory interface on exact-match nuScenes test Actors. Cost and error are row-level 90th percentiles over retained V6.7 outcomes; FS/flip gives event counts. Physical selection is not trajectory authority.

P4 group	Actors	q90 cost $\downarrow$	q90 error $\downarrow$	FS / flip $\downarrow$
selected	5	0.283	1.577	0 / 0
abstained	83	0.102	0.693	7 / 47
geometric harm	23	0.076	0.543	0 / 4

contains 88 test Actors and 39,285 rows (Table 5). P4 selects 5 Actors and abstains on all 23 whose compiled Chamfer is worse than the query; selected-and-harmful is therefore zero. Its 1,781 selected rows contain no false-safe or decision-flip events. Nevertheless, selected q90 motion error is 1.577 m versus 0.693 m for abstained Actors ($2.28\times$), reaching 4.246 m at 3 s. The physical gate protects representation quality but cannot replace the downstream reliability authority. These are retained source outcomes, not a causal effect of repair or an execution of the reliability model.

We also stress the analytic geometry-to-cost bound on 575,596 retained profiles under 6 fixed $\pm 0.05/0.10/0.20$ m shifts. FP64 evaluation has zero violations. At -0.20 m, full-source mean/q99 absolute cost shift is 0.0292/0.1450, while only 0.841% of rows cross signed clearance. P4-selected versus abstained mean shift is 0.00085/0.00456 ($5.38\times$ more stable) with no selected crossings. Selected Actors can therefore be motion-difficult yet geometrically stable, further separating the three axes of repairability, geometry sensitivity, and trajectory uncertainty.

Minimal composed-authority utility. Task-relevant monitoring should propagate prediction error to the queried planning cost [3], rather than treat all prediction errors alike. We therefore form a retained-source 2×2 audit over query/HARP-3D surfaces and no-task/frozen-P346 authority. The 88 exact Actors (32 hazardous) yield 1,228 fixed two-action sets at the P346 artifact’s held-out 2.5 s horizon. Frozen 90% requested reliability authorizes 49.43% of sets, reducing conditional mean actual cost from 0.1497

Table 6. Retained-source 2×2 composition on 88 exact P5 Actors (32 hazardous), all retained. Coverage/cost/unsafe are over fixed two-action sets (all sets for “none”). Equal B0/B1 and B2/B3 action metrics are the intended physical/task non-interference contract, not a causal planning effect.

Arm (surface / authority)	CD $\downarrow$	cov. $\uparrow$	cost $\downarrow$	unsafe $\downarrow$
B0 query / none	0.229	100.00	0.1497	16.45
B1 HARP-3D / none	0.223	100.00	0.1497	16.45
B2 query / P346	0.229	49.43	0.0279	1.15
B3 HARP-3D / P346	0.223	49.43	0.0279	1.15

to 0.0279 (81.37%) and unsafe-set rate from 16.45% to 1.15%. Independently, selective HARP-3D lowers mean surface Chamfer from 0.229 to 0.223 m with all Actors and hazards retained (Table 6). Physical branches have identical action outcomes by construction: this validates interface composition, not closed-loop benefit from repair.

Qualitative protocol. Main-paper and supplement scenes are frozen from metadata and registered failure strata before rendering method outputs. Panels show RGB, LiDAR, physical surface, ray evidence, compiler action, and the repaired world. We report all frozen cases, including failures, rather than selecting only visually favorable examples.

6. Limitations

The original AV2 attribution is a target-nearest proxy. A consumed correction and prospective disjoint AV2 read establish the operator gap, but the latter transfers a metric and provenance—not a learned repair policy. Point deletion certifies non-increasing early returns, not hit retention, Chamfer, collision freedom, or policy safety. Sparse sensing leaves unknown space; tracked cuboids may be wrong; only 64.0% of Actors add no observed-ray violation. Abstention can sacrifice geometry or hazardous-Actor coverage. The 29-Actor source fit, 8.9%→75.6% coverage shift, rank reversal, and 49.5% AV2 sensor-opportunity sensitivity rule out cross-domain exchangeability. Reliability alignment has five selected test Actors, while composition reuses source outcomes without vehicle dynamics. Static topology, deformable Actors, measured sensor-error sets, closed-loop dynamics, and appearance-dependent hazard remain open; HARP-3D is not a real-road safety guarantee.

7. Conclusion

HARP-3D separates validity, hazard, and task authority. Literal audits expose proxy undercounting and deletion monotonicity—not road-safety guarantees.

References

- [1] Julian Bitterwolf, Alexander Meinke, and Matthias Hein. Certifiably adversarially robust detection of out-

- of-distribution data. In *Advances in Neural Information Processing Systems*, 2020. 2
- [2] Holger Caesar, Varun Bankiti, Alex H. Lang, Sourabh Vora, Venice Erin Liong, Qiang Xu, Anush Krishnan, Yu Pan, Giancarlo Baldan, and Oscar Beijbom. nuscenes: A multi-modal dataset for autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020. 2
- [3] Alec Farid, Sushant Veer, Boris Ivanovic, Karen Leung, and Marco Pavone. Task-relevant failure detection for trajectory predictors in autonomous vehicles. In *Proceedings of the 6th Conference on Robot Learning*, pages 1959–1969, 2023. 8
- [4] Yonatan Geifman and Ran El-Yaniv. Selectivenet: A deep neural network with an integrated reject option. In *Proceedings of the 36th International Conference on Machine Learning*, pages 2151–2159, 2019. 2
- [5] Peiyun Hu, Jason Ziglar, David Held, and Deva Ramanan. What you see is what you get: Exploiting visibility for 3d object detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11001–11009, 2020. 2
- [6] Jonas Kälble, Sascha Wirges, Maxim Tatarchenko, and Eddy Ilg. Accurate training data for occupancy map prediction in automated driving using evidence theory. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5281–5290, 2024. 2, 7
- [7] Jonas Kälble, Sascha Wirges, Maxim Tatarchenko, and Eddy Ilg. Evocc: Accurate semantic occupancy for automated driving using evidence theory. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 27467–27476, 2025. 2
- [8] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkuehler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics*, 42 (4), 2023. 2
- [9] Hyeonseong Kim, Yoonsu Kang, Changgyoon Oh, and Kuk-Jin Yoon. Single domain generalization for lidar semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17587–17598, 2023. 2, 6
- [10] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, 2017. 2
- [11] David Madras, Toni Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and accuracy by learning to defer. In *Advances in Neural Information Processing Systems*, 2018. 2
- [12] Julian Ost, Fahim Mannan, Nils Thuerey, Julian Knodt, and Felix Heide. Neural scene graphs for dynamic scenes. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021. 2
- [13] Yaniv Romano, Evan Patterson, and Emmanuel J. Candès. Conformalized quantile regression. In *Advances in Neural Information Processing Systems*, 2019. 2
- [14] Jules Sanchez, Jean-Emmanuel Deschaud, and François Goulette. Domain generalization of 3d semantic segmentation in autonomous driving. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 18077–18087, 2023. 2
- [15] Nikolay Savinov, Christian Hane, Lubor Ladicky, and Marc Pollefeys. Semantic 3d reconstruction with continuous regularization and ray potentials using a visibility consistency constraint. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 5460–5469, 2016. 2
- [16] Abhin Shah, Yuheng Bu, Joshua K. Lee, Subhro Das, Rameswar Panda, Prasanna Sattigeri, and Gregory W. Wornell. Selective regression under fairness criteria. In *Proceedings of the 39th International Conference on Machine Learning*, pages 19598–19615, 2022. 2
- [17] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 3319–3328, 2017. 6
- [18] Adam Tonderski, Carl Lindström, Georg Hess, William Ljungbergh, Lennart Svensson, and Christoffer Petersson. Neurad: Neural rendering for autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. 4
- [19] Benjamin Wilson, William Qi, Tanmay Agarwal, John Lambert, Jagjeet Singh, Siddhesh Khandelwal, Bowen Pan, Ratnesh Kumar, Andrew Hartnett, Jhony Kaesemodel Pontes, Deva Ramanan, Peter Carr, and James Hays. Argoverse 2: Next generation datasets for self-driving perception and forecasting. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021. 2
- [20] Han Zhou, Jordy Van Landeghem, Teodora Popordanoska, and Matthew B. Blaschko. A novel characterization of the population area under the risk coverage curve (aurc) and rates of finite sample estimators. In *Proceedings of the 42nd International Conference on Machine Learning*, pages 79226–79253, 2025. 7